

ESGALDNS:
NAMECOIN-BASED ANONYMOUS DOMAIN NAME SERVICE
FOR TOR HIDDEN SERVICES

by

Jesse Victors

A thesis submitted in partial fulfillment
of the requirements for the degree

of

MASTER OF SCIENCE

in

Computer Science

Approved:

Dr. Ming Li
Major Professor

Dr. Nicholas Flann
Committee Member

Dr. Daniel Watson
Committee Member

Dr. Mark R. McLellan
Vice President for Research and
Dean of the School of Graduate Studies

UTAH STATE UNIVERSITY
Logan, Utah

2014

Copyright © Jesse Victors 2014

All Rights Reserved

ABSTRACT

EsgalDNS:
Namecoin-based Anonymous Domain Name Service
for Tor Hidden Services

by

Jesse Victors, Master of Science
Utah State University, 2014

Major Professor: Dr. Ming Li
Department: Computer Science

Web applications are increasingly moving business and processing logic from the server to the browser. Traditional, multiple-page request/response applications are quickly being replaced by single-page applications where complex application logic is downloaded on the initial page load and data is then subsequently fetched asynchronously via the browser's native XMLHttpRequest (XHR) object.

These new generation web applications are called Rich Web Applications (RWA). Frameworks such as the Google Web Toolkit (GWT), and JavaScript model-view-controller (MVC) frameworks such as Backbone.js are facilitating this move. With this migration, testing frameworks need to follow the logic by moving analysis and test generation from the server to the client. One problem hindering the movement of testing in this domain is the adoption of semantic URLs. This paper introduces a novel approach to systematically identify variables in semantic URLs and use them as part of the test generation process.

Using a sample RWA seeded with various JavaScript faults, I demonstrate in this thesis, as an empirical study, that combinatorial testing algorithms and reduction strategies also apply to new RWAs.

PUBLIC ABSTRACT

CHAD M. MAUGHAN

Rich Web Applications (RWA) that are data driven and feature responsive user interfaces are rapidly growing in popularity. Popular sites such as Twitter, Pandora, and Angry Birds (browser version) are all examples of popular RWAs. These RWAs are more complex and are developed differently than many previous web sites. More of the processing power needed to run these applications is performed on the client machine, not the server. Due to this development strategy, testing tools need new techniques to identify web application variables, capture errors, and identify problems. In this thesis, I introduce novel techniques to identify variables in RWA semantic URLs and automatically generate tests for RWAs using a form of testing called combinatorial testing.

CONTENTS

	Page
ABSTRACT	iii
PUBLIC ABSTRACT	v
LIST OF TABLES	vii
LIST OF FIGURES	viii
CHAPTER	
1 INTRODUCTION	1
1.1 Tor	1
1.2 Bitcoin	6
1.3 Namecoin	9
1.4 DNS	9
2 BACKGROUND	10
2.1 JavaScript Background	10
2.2 Traditional URL Formats	11
2.3 Semantic URL Formats	11
2.4 Variable Detection Algorithm for Semantic URLs	12
2.5 Algorithm Description	13
3 EXPERIMENTS	16
3.1 Research Questions	16
3.2 Sample Application	16
3.3 JavaScript Faults	17
3.4 Testing Technologies	20
3.5 Navigation Graph	22
3.6 Abstract URLs	22
3.7 Capturing Errors	25
3.8 Combinatorial Testing	26
4 RESULTS	27
4.1 Size Impact of Abstract URL Test Suite	27
4.2 Effectiveness of Abstract URL Test Suite	28
4.3 Size Impact of Combinatorial Coverage	28
4.4 Effectiveness of Combinatorial Coverage	30
5 CONCLUSIONS	31
REFERENCES	34

LIST OF TABLES

Table	Page
3.1 Sample Rich Web Application Information.	17
3.2 Sample Rich Web Application Seeded Errors.	21
4.1 Size of Test Suites by Strategy.	28
4.2 Effectiveness of Single Coverage With Abstract URL Reduction.	28
4.3 Size of Test Suites by Strategy.	28
4.4 Sample Combinatorial Tests.	29
4.5 Effectiveness of Combinatorial Coverage.	30
5.1 Environment Variables.	32

LIST OF FIGURES

Figure	Page
1.1 The user’s software first downloads a directory list of Tor relays. This information is later used to construct a circuit through the network. [?]	2
1.2 A Tor circuit is incrementally constructed using layers of encryption. Each node has limited visibility, and no individual node knows the whole circuit.	3
1.3 Construction of a Tor circuit.	4
1.4 A Tor circuit is changed periodically, essentially providing a new identity to the end user. Although a different guard node is used here, in practice the choice of entry point is preserved for extended periods of time. [?]	4
1.5 Alice uses the encrypted cookie to tell Bob the secret and to switch to <i>RP</i> . [?]	6
1.6 Alice and Bob can now anonymously communicate over their Tor circuit. The communication travels through six Tor nodes: three established by Alice and three by Bob. [?]	6
1.7 In the possibility that multiple nodes solve the proof-of-work and generate a new block simultaneously, the block becomes orphaned, the transactions recycled, and the blockchain follows the longest path from the genesis node to the latest block.	8
1.8 Three traditional Bitcoin transactions. Each transaction contains the public of the recipient, the ECDSA digital signature of the transaction from the sender, and the hash of the originating transaction. In practice transactions can contain multiple input and outputs.	8
2.1 Algorithm for variable identification of semantic URLs	14
2.2 A visual representation of branching complexity in an application structure.	15
3.1 Sample rich web application screenshot.	17
3.2 Depth first crawling of a rich web application	23
3.3 An example JavaScript error in the Firebug console.	25
3.4 Browser exception catching	26
4.1 Hierarchical variable HTML form example	29

CHAPTER 1

INTRODUCTION

1.1 Tor

The Tor network is a second-generation onion routing system that aims to provide anonymity, privacy, and Internet censorship protection to its users. Tor routes encrypted TCP/IP traffic through a worldwide volunteer-run network of over six thousand relays. Tor's encryption, authentication, and routing protocols are designed to make it infeasible for any adversary to identify an end user or reveal their traffic. Despite attempts by agencies and governments to block, tap, or crack the Tor network, Tor remains one of the most popular and secure tools to use against network surveillance, traffic analysis, and information censorship.

1.1.1 Design

Tor provides an anonymity and privacy layer by relaying all end-user TCP traffic through a series of relays on the Tor network. Typically this route consists of a carefully-constructed three-hop path known as a *circuit*, which changes over time. These nodes in the circuit are commonly referred to as *guard node*, *middle relay*, and the *exit node*, respectively. Only the first node can determine the origin of TCP traffic into Tor, and only the exit node can see the destination of traffic out of Tor. The middle router is unable to determine either. Furthermore, each node is only aware of the machines it talks to, so only the client knows the identity of all three nodes used in its circuit. Tor's architecture is designed to make it exceptionally difficult for a well-resourced adversary to uncover the identity of the end-user and their network activities, even if nodes are compromised. [?]

1.1.2 Routing

In traditional Internet connections, the client communicates directly with the server. In this model, an eavesdropper can often reveal both the identity of the end user and their activities. Direct encrypted connections do not hide IP headers, which expose source and destination addresses and the size of the payload. In the face of adversaries with sophisticated traffic analysis tools, such information can be very revealing for someone who wishes to hide their activities altogether.

Tor combats this by routing end user traffic through a randomized circuit through the network of relays. The Tor client software first queries a trusted directory server or a relay mirroring the directory. This directory contains of list of IPs, ports, public keys, and other information about all nodes in the Tor network. [?] This query is illustrated in Figure 1.

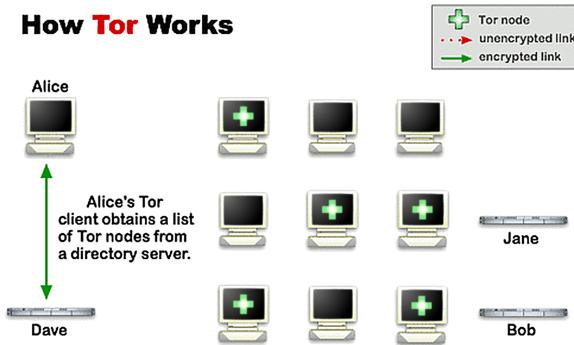


Figure 1.1: The user’s software first downloads a directory list of Tor relays. This information is later used to construct a circuit through the network. [?]

Next, the Tor client chooses three unique and geographically diverse nodes to use. It then builds and extends the circuit one node at at time, negotiating respective HTTPS connections with each node in turn. No single relay knows the complete path, and each relay can only decrypt its layer of decryption. In this way, data is encrypted multiple times and then is decrypted in an onion-like fashion as it passes through the circuit.

The client first establishes a TLS connection with the first relay, R_1 , using the relay’s

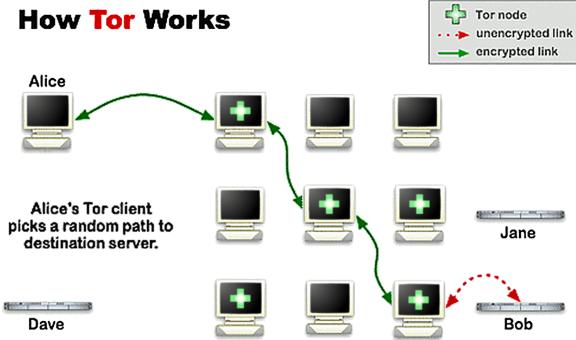


Figure 1.2: A Tor circuit is incrementally constructed using layers of encryption. Each node has limited visibility, and no individual node knows the whole circuit.

public key. The client then performs a Diffie-Hellman-Merkle key exchange to negotiate K_1 which is then used to generate two symmetric session keys: a forward key $K_{1,F}$ and a backwards key $K_{1,B}$. $K_{1,F}$ is used to encrypt all communication from the client to R_1 and $K_{1,B}$ is used for all replies from R_1 to the client. These keys are used conjunction with the symmetric cipher suite negotiated during the TLS handshake, thus forming an encrypted tunnel with perfect forward secrecy. Once this one-hop circuit has been created, the client then sends R_1 the RELAY_EXTEND command, the address of R_2 , and the client's half of the Diffie-Hellman-Merkle protocol ($g^a mod b$) using $K_{1,F}$. R_1 performs a TLS handshake with R_2 and uses R_2 's public key to send $g^a mod b$ to R_2 , who replies with his half of the handshake and a hash of K_2 . R_1 then forwards this to the client under $R_{1,B}$ with the RELAY_EXTENDED command to notify the client. The client generates $K_{1,F}$ and $K_{1,B}$ from K_2 , and repeats the process for R_3 , [?] as shown in Figure 3. The TLS/IP connections remain open, so the returned information travels back up the circuit to the end user.

Following the complete establishment of a circuit, the Tor client software then offers a Secure Sockets (SOCKS) interface on localhost which multiplexes TCP traffic through Tor. At the application layer, this data is packed and padded into equally-sized Tor *cells*, transmission units of 512 bytes. As each relay sees no more than one hop in the circuit, in theory neither an eavesdropper nor a compromised relay can link the connection's source,

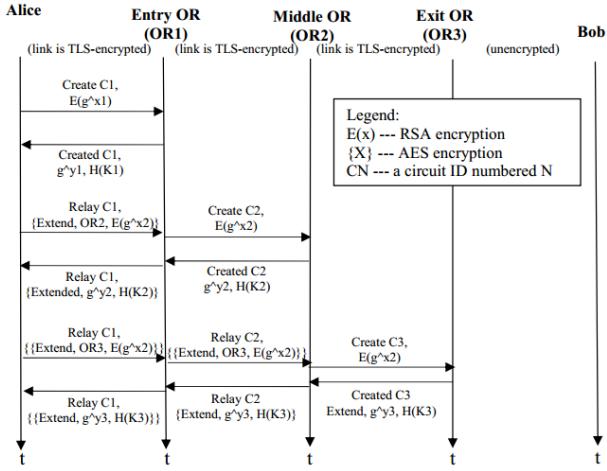


Figure 1.3: Construction of a Tor circuit.

destination, and content. Tor further obfuscates user traffic by changing the circuit path every ten minutes, [?] as shown in Figure 4. A new circuit can also be requested manually by the user.

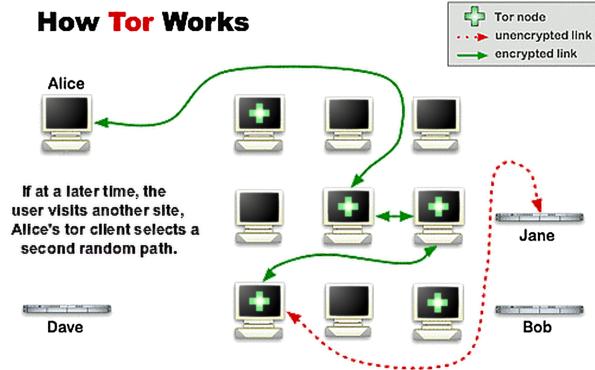


Figure 1.4: A Tor circuit is changed periodically, essentially providing a new identity to the end user. Although a different guard node is used here, in practice the choice of entry point is preserved for extended periods of time. [?]

An encrypted connection is often established with the other client or web server, depending on whether or not the recipient supports encryption. If this is the case, even the exit node cannot see the traffic in cleartext. An outsider is therefore faced with up

to four layers of TLS encryption: $K_{1,F}(K_{2,F}(K_{3,F}(K_{server}(\text{client request}))))$ and likewise $K_{1,B}(K_{2,B}(K_{3,B}(K_{server}(\text{server reply}))))$ for the returning traffic. This makes traffic analysis and cryptographic attacks very difficult.

Tor users typically use the Tor Browser Bundle, (TBB) a custom build of Mozilla Firefox with a focus on security and privacy. The TBB anonymizes and provides privacy to the user in many ways. These include blockin

The TBB not only provides special handling of client-side scripts such as Javascript, but also offers the HTTPS Everywhere extension, which uses regular expressions to rewrite HTTP web requests into HTTPS whenever possible. Thus, if the web server is capable of handling SSL or TLS connections, HTTP communications will be encrypted to them. If this is the case, the TBB performs a TLS handshake with the web server, but the exchange happens through the Tor circuit. This provides the final layer of encryption to the outside.

1.1.3 Hidden Services

While the majority of Tor's usage is for traditional access to the Internet, Tor's routing scheme also supports anonymous services, such as websites or chatrooms. These are a part of the Deep Web and cannot be normally contacted outside of Tor. Unlike the Clearnet, Tor does not contain a traditional DNS system for its websites; instead, every hidden service has a public and private RSA key, and domains are a truncated SHA-1 hash of its public key. This means that domain names are distributed and correlate directly to the identity of a hidden service, allowing anyone to verify the authenticity of the service server, akin to SSL certificates on the Clearnet. Tor hidden services allow a client, Alice, and a hidden service, Bob, to communicate anonymously. [?]

In advance, Bob builds Tor circuits to several random relays and enables them to act as *introduction points* by giving them his public key, B_K . He then uploads this information to a distributed hashtable inside the Tor network, signing the result. Alice queries this hashtable, finds B_K and his introduction points, and builds a Tor circuit to one of them, IP_1 . Simultaneously, she also builds a circuit to another relay, RP , which she enables as a rendezvous point by telling it a one-time secret, S .

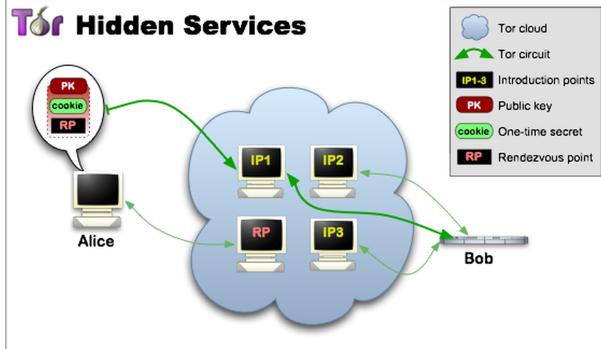


Figure 1.5: Alice uses the encrypted cookie to tell Bob the secret and to switch to *RP*. [?]

She then sends to IP_1 an cookie encrypted with B_K , containing RP and S . Bob decrypts this message, builds a circuit to RP , and tells it S_1 , enabling Alice and Bob to communicate.

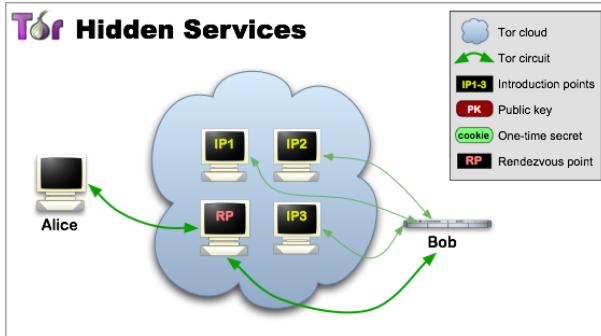


Figure 1.6: Alice and Bob can now anonymously communicate over their Tor circuit. The communication travels through six Tor nodes: three established by Alice and three by Bob. [?]

1.2 Bitcoin

Bitcoin is a decentralized digital cryptocurrency, created by pseudonymous developer Satoshi Nakamoto in 2008. Ownership of Bitcoins consists of holding a private ECDSA key, and a transfer is simply a transmission of Bitcoins from one key to another. All transactions are recorded on a public ledger, called a blockchain, a data structure whose

integrity is verified through computational power. Bitcoins are generated computationally at a fixed rate by *miners*, who also secure the blockchain. Although Bitcoin received limited attention in the first two years of its life, it has since grown significantly since then, with approximately 70,000 daily transactions as of the time of this writing. Bitcoin's growth has led to the creation of many alternative cryptocurrencies, and its popularity has influenced financial discussions, legal controversy, and the prices of electronics worldwide.

1.2.1 Architecture

A blockchain is data structure fundamental to Bitcoin, and crucial for its functionality. As a distributed decentralized system, this public ledger is Nakamoto's answer to ensuring agreement of fundamental data across all involved parties. The blockchain is a novel structure, and its structure guarantees integrity, chronological ordering of transactions, and prevention of double-spending of Bitcoins. The blockchain consists of blocks of data that are held together by proof-of-work, a cryptographic puzzle whose solution is provably hard to find but trivial to verify. Bitcoin's proof-of-work is based on Adam Back's Hashcash scheme: that is, find a nonce such that the hash of this nonce meets a certain requirement. In Bitcoin's case this is stated as finding a nonce that when passed through two rounds of SHA-256 (SHA256²) produces a value less than or equal to a target T . This requires a party to perform on average $\frac{1}{Pr[H \leq T]} = \frac{2^{256}}{T}$ amount of computations, but it is easy to verify that $\text{SHA256}^2(\text{msg}||n) \leq T$. Nodes in the Bitcoin network collectively agree to use the blockchain with the highest accumulation of computational effort, so an adversary seeking to modify the structure would need to recompute the proof-of-work for all previous blocks as well as out-perform the network, which is infeasible. [?]

Each block in the blockchain consists of a header and a payload. The header contains a hash of the previous block's header, the root hash of the Merkle tree built from the transactions in this block, a timestamp, a target T , and a nonce. The target T changes in response to speed at which the proof-of-work was solved such that blocks are added to the blockchain at a relatively fixed rate. The block payload consists of a list of transactions. The root node of the Merkle tree ensures the integrity of the transaction vector: verifying

that a given transaction is contained in the tree takes $\log(n)$ hashes, and a Merkle tree can be built $n * \log(n)$ time, ensuring that all transactions are accounted for. The hash of the previous block in the header ensures that blocks are ordered chronologically, and the Merkle root hash ensures that the transactions contained in each block are ordered chronologically as well. The SHA256² proof-of-work provides integrity of the data structure, and secp256k1 ECDSA key are used to prove ownership of coins. [?]

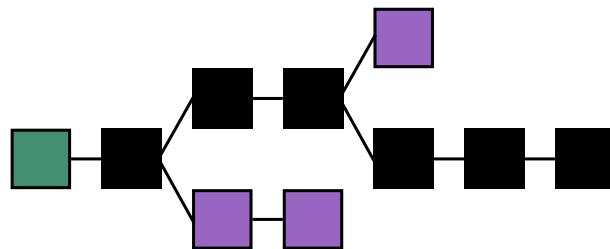


Figure 1.7: In the possibility that multiple nodes solve the proof-of-work and generate a new block simultaneously, the block becomes orphaned, the transactions recycled, and the blockchain follows the longest path from the genesis node to the latest block.

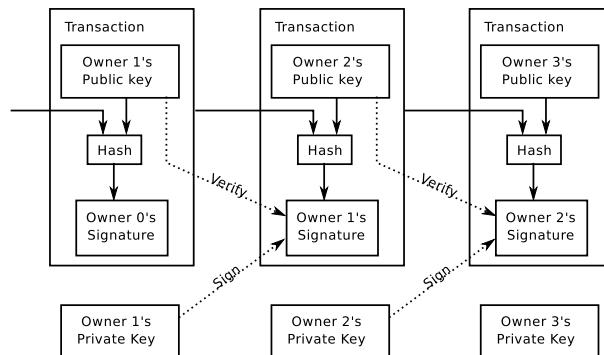


Figure 1.8: Three traditional Bitcoin transactions. Each transaction contains the public key of the recipient, the ECDSA digital signature of the transaction from the sender, and the hash of the originating transaction. In practice transactions can contain multiple input and outputs.

In this way, the digital signatures and proof-of-work in the blockchain can be tracked forwards indefinitely.

1.3 Namecoin

1.4 DNS

1.4.1 Introduction

1.4.2 Architecture

CHAPTER 2

BACKGROUND

This chapter is divided into four sections. First, I provide some background on JavaScript and the characteristics of the language that make it challenging to test. Second, I explain a traditional URL format, describing the relative ease of programmatically discovering variables. Third, I explain the more recent semantic URL format and some of the difficulties introduced with systemic variable identification. Fourth, I introduce and describe an algorithm for statistically analyzing an application’s URL structure to identify variables.

2.1 JavaScript Background

Due to its distribution in all browsers, JavaScript is the primary language for Rich Web Applications (RWA). It plays a central role in RWAs by updating styles and modifying application structure markup (i.e., the Document Object Model or DOM). It also is the key data transport mechanism for interacting with the server via its XMLHttpRequest (XHR) object. As a language, it is dynamic, weakly typed, interpreted, and prototype-based. With its features such as first-class functions, it supports multiple development paradigms, including: object-oriented, imperative, and functional. JavaScript can either be loaded by the browser as static code or it can be dynamically created at runtime using an “`eval()`” function.

Two traits make JavaScript particularly vulnerable to faults. First, as a dynamically typed, interpreted language it does not benefit as much as strongly typed, compiled languages from static code analysis. Lint-like tools exist for JavaScript but they typically focus on syntax. Second, JavaScript is designed to respond to timers and browser events such as clicks, hovering, or key presses. This makes it difficult to test event driven code in a development environment as it requires simulating those events to test.

New JavaScript frameworks such as Backbone.js have helped with the consistency of JavaScript code being developed. They have helped isolate faults to mostly appearance related exceptions. In the next chapter, I describe in detail the types of web application and JavaScript faults in the context of a sample application and how those faults are identified with the developed testing strategies.

2.2 Traditional URL Formats

Rich web applications, along with a migration from server-side to client-side JavaScript based frameworks, have generally adopted a different URL structure from the traditional format defined in RFC 1738 [1]. A traditional URL structure defined in the RFC 1738 standard, follows the format:

```
http://<host>:<port>/<path>?<searchpart>
```

The <searchpart> of a traditional URL format is composed of a series of key-value pairs that start after the question mark. Keys and values are separated by an equals sign and the key-value pairs are separated by an ampersand. An example of a <searchpart>, more frequently called a query string, following the question mark of a URL is as follows:

```
key1=value1&key2=value2
```

Traditional URL structures, as defined in RFC 1738, with their key-value pairs are very easy to programmatically identify and parse. This well structured format allows a testing solution to quickly determine variable names and values for combinatorial testing.

2.3 Semantic URL Formats

Rich web applications have departed from this traditional <searchpart> format in favor of a more user friendly, more permanent format [2]. This new format is often referred to as semantic or clean URLs. One goal of semantic URLs is to provide an immediate, human

understandable format of resources. Instead of clearly identifying variables and their values as key-value pairs after a question mark in the <searchpart> section of a URL, rich web applications typically combine their variables and values as part of the <path> portion of the URL. While easier for a human to understand, this structure makes it difficult to programmatically discover variables for combinatorial testing.

As an example of a semantic URL, imagine a rich web application that provides census information on states, counties, and cities over a certain population of the United States. A semantic URL for this application might look as follows:

```
http://example.com/#/state/utah/county/cache/city/logan
```

This example is a well structured semantic URL. Each variable is preceded by a descriptive key (e.g., “state” precedes “ut” and “county” precedes “cache”). Unfortunately, no standards exist for semantic URL structures, and many rich web applications do not always find it practical to follow this industry best practice of associated keys and values for variables. This application could have just as easily been developed with a URL structure:

```
http://example.com/#/utah/cache/logan
```

Due to the variety of formats used in semantic URLs and the absence of an industry standard, identifying variables for combinatorial testing can be difficult. In the next section I propose an algorithm for variable identification regardless of the example structures shown above.

2.4 Variable Detection Algorithm for Semantic URLs

In order to address the complexities of identifying variables in semantic URLs, I propose the following algorithm that processes all URLs from a rich web application and identifies the location of variables in a URL structure.

The algorithm processes each URL of the rich web application and then subsequently processes each portion of the URL <path>. It then combines them into a common hierarchical graph to allow analysis of the application structure. The key to identifying the location of variables is by calculating branching complexity as each URL is processed. This branching complexity is measured by creating a graph node at each URL <path> portion (e.g., each part of the URL path separated by a “/” character). As each node is created, the Strahler number is calculated by looking at all the previous child nodes [3]. Identifying the degree of branching at each level of the URL path allows for analysis of the branching complexity for each abstract portion of the URLs of the application.

The algorithm is listed in Figure 2.1 followed in the next section by a detailed description.

2.5 Algorithm Description

For brevity, the algorithm, as listed, assumes the possession of a list every possible URL available in the application. The sample rich web application discussed in this section has approximately 200,000 different semantic URLs. A better implementation of the algorithm would be to crawl the site in a depth-first, non-deterministic fashion recording and processing each URL (i.e., applying this algorithm in-line as the site is crawled).

The algorithm proceeds to lob off each outer part of the URL path until the position variable listed on line 103 is less than the fragment identifier character position in the URL.

Starting with the example URL listed previously, the algorithm values are processed in the following list.

1. `http://example.com/#/state/ut/county/cache/city/logan.`
2. `http://example.com/#/state/ut/county/cache/city`
3. `http://example.com/#/state/ut/county/cache`
4. `http://example.com/#/state/ut/county`
5. `http://example.com/#/state/ut/`

```

100  foreach(url in urls) {
101      Node previousNode = retrieveOrCreateNode(url);
102      int position = url.length();
103      String shortenedUrl;
104      while(position > fragmentIdentifier) {
105          position = url.lastIndexOf("/");
106          if(position > 0) {
107              shortenedUrl = url.substring(0, position);
108              Node node = retrieveOrCreateNode(shortenedUrl);
109              int branchingComplexity = calculateBranchingComplexity(node);
110              node.setProperty(BRANCHING.COMPLEXITY, branchingComplexity);
111              previousNode.createRelationshipTo(node,
112                                              RelationshipTypes.CHILD_OF);
113              previousNode = node;
114          }
115          else {
116              break;
117          }
118      }
119      Node last = retrieveOrCreateNode(shortenedUrl);
120  }

```

Figure 2.1: Algorithm for variable identification of semantic URLs

6. <http://example.com/#/state>

7. <http://example.com/#>

As we're focused on the <path> portion of the URL, the algorithm stops processing when the position index of the last instance of the path separator "/" is less than the position index of the fragment identifier (the "#" character). Continuing on line 108, the algorithm then either retrieves (if it already exists) or creates (if the remaining URL has not been processed) a new node that represents that portion of the URL. The algorithm works its way backwards from full URL to just the host to allow for the easy calculation of the Strahler number used to identify branching complexity in line 109. On line 111 the

algorithm creates the edge between the current and the previous node. It then assigns the current node to the previous node so it can continue processing if needed.

An additional benefit of having the application structure in a common graph is it makes it easy to visually identify branching complexity as in Figure 2.2. Dark areas represent extensive branching that is associated with variables in semantic URLs.

Figure 2.2: A visual representation of branching complexity in an application structure.

CHAPTER 3

EXPERIMENTS

3.1 Research Questions

There are four main questions of focus that the following experiments attempt to answer:

1. What is the impact on the size of the test suite using abstract URLs instead of an exhaustive enumeration of every variable combination?
2. What is the fault finding effectiveness of testing with abstract URLs versus exhaustive testing?
3. What is the impact on the size of the test suite using combinatorial coverage compared to single coverage?
4. What is the fault finding effectiveness of the combinatorial coverage compared with single coverage?

3.2 Sample Application

All experiments were conducted on a sample Rich Web Application written by the author and is available at <http://chadmaughan.com/thesis>. The application provides United States census information for 1990, 2000, and preliminary numbers from 2005 for each state, county, and city with a population over 25,000 residents. The application is designed for mobile devices with larger touch points. The sample application uses the following technologies:

1. Backbone.js: A client-side, JavaScript MVC framework that gives structure to rich web applications. Backbone.js allows you to bind custom events to models, and route URL fragments to JavaScript functions.
2. jQuery Mobile: A unified, HTML5-based cross-platform user interface system for mobile devices. It focuses on semantic markup that is easily themeable.
3. Spring MVC: A module of the popular Spring Framework for Java that allows for easy implementation of server-side controllers for building REST APIs.
4. Apache Derby DB: An open source, embedded relational database implemented in Java.

Source code for both the application and the fault finding, testing code is also available at <http://code.chadmaughan.com/thesis>. Metrics about the sample application are included in Table 3.1. A sample screenshot of the application is provided in Figure 3.1.

Table 3.1: Sample Rich Web Application Information.

Number of Application States	199,484
Number of Files	328
JavaScript Lines of Code	605
Java Number of Classes	21
Java Lines of Code	1,091
Seeded Faults	73

Figure 3.1: Sample rich web application screenshot.

3.3 JavaScript Faults

The sample application is seeded with 73 faults. A detailed list of those exceptions and the categories they belong to are listed in Table 3.2.

On fault categories, Sampath et al. described five web application fault categories [4], of which all are used in the sample application.

1. Data store faults: Faults in the application code that manipulates data in any kind of data store. This category of faults also applies to data that is incorrectly persisted in the data store. There are a number of seeded data store faults in the sample application.
2. Logic faults: Faults in the application code that implements business logic and control flow. An example of this is an error in the page transition with jQuery Mobile.
3. Form faults: Faults in the application code that controls, modifies and displays name-value pairs in forms. The sample application does not submit any data to the server but does use dynamically updated forms to direct to different application states.
4. Appearance Faults: faults in the application code that controls the way in which a web page is displayed. With the use of modern JavaScript based templating solutions, such as Mustache.js and Underscore.js, appearance faults typically cause a template to not be rendered. The sample application demonstrates this type of error when trying to load the search page.
5. Link faults: Faults in the application code that changes the page pointed to by an URL.

Guo and Sampath [5] add a sixth category, compatibility faults or “faults in application code that ensures that the web application complies with different browsers, versions of browsers and other client environments.” With the rapid introduction of new HTML5 APIs (i.e., Websockets and Webstorage) and the varying speed of adoption among browser creators, compatibility faults will play an increasingly important role in client-side testing.

Guo also expands the logic faults category to include seven sub-categories. The sub-categories of logic faults are:

1. Browser interaction faults: Faults in the application code that control the web browser, such as code that disables the “back” button on the browser, or code that is affected by user-defined browser settings, such as disabled cookies. Browser manipulation is discouraged as it alters expected application behavior. Other browser settings, such as disabling JavaScript, would render the application useless. As such, the sample application does not use this fault sub-category.
2. Session faults: Faults in the application code that deal with maintaining state of application or other session-based operations such as using sessions to save and data entered into a form and display the data after the sessions has been validated. As most rich web applications are stateless to avoid the overhead of session management and allow for greater scalability, the sample application does not use this sub-category.
3. Paging faults: Faults in the application code that deals with paging when displaying large amounts of data on the screen. While applicable to rich web applications, the sample application does not use this sub-category.
4. Server-side parsing faults: Faults in the application code that deal with server-side parsing of HTML, XML, and JavaScript tags. This sub-category does not apply to rich web applications as HTML used in a RWA is typically a minimal page used only as an entry point. “client-side parsing faults” are more applicable to rich web applications. The sample application has a number JavaScript syntax related faults. These errors are described below in Table 3.2.
5. Encoding/decoding faults: Faults in the application code that encodes or decodes characters for transmission, storage, and display
6. Locale faults: Faults that exist in application code that sets or gets locale-specific information, such as date format or language. Not used in the sample application.
7. Other: Other logic faults that do not belong to any of the above sub categories.

More recently, in addition to Sampath and Guo, Ocariza et al. [6] defines five categories specifically for JavaScript exceptions. He found that JavaScript exceptions tend to be much more defined than other applications and fall into well-defined categories. In fact, 94% of all errors studied from the Alexa top 100 list fall into five categories, namely:

1. Permission Denied: Faults in the application code that attempt to access JavaScript components from another domain violating the same-origin policy. The sample application has a seeded fault where it tries to load recent search data from Twitter and violates the same-origin policy.
2. Null Exception faults: Faults in the application code where a property or method is accessed via a null object. The sample application attempts to add a CSS class name to an element that is null.
3. Undefined Symbol faults: Faults in the application code where a function or variable is accessed that has not been previously defined.
4. Syntax Error faults: Faults in the application code where interpreted code, such as in an eval() function, has the wrong syntax.
5. Miscellaneous faults: Faults in the application code that apply specifically to a single site.

These five categories align more closely with the JavaScript Error object and its six other core errors: EvalError (Syntax Error faults), RangeError, ReferenceError (Undefined Symbol faults), SyntaxError (Syntax Error faults), TypeError, and URIError [7].

3.4 Testing Technologies

In addition to the sample Rich Web Application, the code used to identify the seeded faults rely on some key technologies. These technologies are listed below:

Table 3.2: Sample Rich Web Application Seeded Errors.

Location	Description	Category	Count
index.html:39	console.log(variable) - variable is not defined.	ReferenceError, Undefined	1
main.js:54	SearchPageView template is not included as a source.	ReferenceError, Appearance	1
geochart:438	County Map doesn't exist on County Details page. "Container is not defined" error.	Appearance	1 (3,144)
main.js:49	Syntax Error on eval() - eval("window.print();")	SyntaxError	1
main.js:37	Missing script file - not-there.js"	Link, Other, Misc	1
jquery-1.7.1.min.js:4	Multiple words in the name (9 states plus Washington DC) doesn't have a flag to display		10
jquery-1.7.1.min.js:4	Multiple words in the name (9 states plus Washington DC) doesn't have a seal to display	Appearance	10
CityController.java:39	Cities with periods in their name/code, return HTTP 500 (i.e., St. George, UT).	Data Store	13
index.html:87	Type Error 'bad type' for all Washington State cities	TypeError, Misc	37
twitter.js:4	City Twitter Search feed	Permission	1 (1,267)
census-table-view.js:17	Adds a CSS class to a null element	Null Exception, TypeError	1

1. Neo4j: A powerful graph database with a rich API that was used to both systematically identify variables in semantic URLs and store a full navigation graph of the rich web application for test traversal retrieval.
2. BrowserMob: An embedded proxy software that allows the interception of HTTP requests from the client and HTTP responses from the server. This allows for the identification of network errors and for the modification of the server responses to assist in identifying client side JavaScript errors.
3. Selenium WebDriver: Allows for the crawling and interaction with the application in the various browsers “as the user would,” giving a more accurate result while testing.

3.5 Navigation Graph

A navigation graph is a visual representation of an application [8]. A navigation graph of a rich web application is a visual representation of how different states of the web application are related to one another. A traditional web application navigation graph typically would have a single node for each HTML page or a single node for each rendering of a server-side script. An edge on a traditional web application represents an HTML anchor tag. As rich web applications typically have a single or very limited number of HTML pages acting as entry points, a node in the navigation graph of a rich web application represents a single state of the application. An edge in a navigation graph for a rich web application represents a transition from one state to another. Like in a traditional web application, this is also typically done through an HTML anchor tag.

Building a navigation graph of the rich web applications has many benefits, including understanding the application as a whole, maintenance over the development cycle, and test case generation and the subsequent testing of those tests [9]. To build the graph, and expand it as the application increases in size, a rich web application can be crawled using the Selenium WebDriver in a simple depth first traversal. Figure 3.2 demonstrates the crawling algorithm using Selenium WebDriver.

Also of benefit is the ability to identify key, holistic characteristics about the application as the navigation graph is built. For example, while this experiment does not focus on test prioritization, systematically creating a navigation graph with Selenium WebDriver allows you to easily record and track key information about state transitions in the application. As an example, one may wish to prioritize testing based on the location of links or elements that trigger a change in state. Combining the exact top and left pixel location of that element allows you to prioritize links that are more in the key navigation sections. As another prioritization strategy with a navigation graph, one could also easily prioritize based on states of the application that are most transitioned.

3.6 Abstract URLs

One problem with the depth-first traversal of a rich web application is the time it

```

100 public static void main(String [] args) {
101     driver = new FirefoxDriver();
102     new Crawl("http://example.com");
103 }
104
105 public Crawl(String startingUrl) throws Exception {
106     String newAddress = null;
107     while ((newAddress = queue.poll()) != null)
108         processPage(newAddress);
109 }
110
111 private void processPage(String sourceHref) throws Exception {
112     driver.get(sourceHref);
113     List<WebElement> links = driver.findElements(By.tagName("a"));
114     Node sourceNode = retrieveOrCreateNode(sourceHref);
115     String targetHref;
116
117     for (int i = 0; i < links.size(); i++) {
118         WebElement w = links.get(i);
119         targetHref = w.getAttribute("href");
120         if (isValidUrl(targetHref)) {
121             if(w.isDisplayed()) {
122                 Node targetNode = retrieveOrCreateNode(targetHref);
123                 if(targetNode == null)
124                     targetNode = nodeMapper.mapNode(targetHref, w);
125                 sourceNode.createRelationshipTo(targetNode,
126                     RelationshipTypes.LINKS_TO);
127                 if (!processed.contains(targetHref))
128                     queue.add(targetHref);
129             }
130         }
131     }
132 }
```

Figure 3.2: Depth first crawling of a rich web application

takes to complete a full application crawl. For example, the sample census application discussed previously contains nearly 200,000 different states. Assuming an average page load of around 500 milliseconds, it would still require more than 27 hours to perform a full regression test. While this amount of time to perform an exhaustive regression test is possible, it may not always be practical.

Using variables previously identified through the algorithm described in the background chapter, application states can be stored in the navigation graph as “abstract URLs” [9], or URLs with variable names instead of each possible value. Similar to semantic URL formats discussed earlier, an example of an abstract semantic URL stored in a navigation graph would look as follows:

```
http://example.com/#/state/{var1}/county/{var2}/city/{var3}
```

Storing the abstract URL instead of every URL variable combination significantly reduces the number of tests required. Indeed, it is ideal as most rich web application states reuse small HTML templates for the displaying of particular data points. This means that client-side errors would typically manifest themselves for every variance of a variable. One disadvantage of using abstract URLs is that data specific errors may be missed. For example, the sample application has a fault where any city with a period in its name (i.e., “St. George, UT”) doesn’t render correctly. Unless the random value for the abstract URL variable selected contained a period, this data related fault would be missed. While not discussed in detail in this paper, due to the available resource of the full application structure (from the systemic variable identification), one possibility is to use a sample size confidence level to adequately test a certain size of the data available.

One of the experiments completed was calculating and comparing the amount of time required to crawl a full rich web application, a reduced number of states based on a sample size, and a fully reduced test suite testing only abstract URLs one time. Results of this experiment are discussed in the following chapter.

3.7 Capturing Errors

Another difficulty encountered with testing rich web applications is the ability to identify errors that occur in the browser. JavaScript errors are difficult to systematically intercept and report. This section introduces a novel approach for identifying two different types of exceptions during testing, namely network related and JavaScript browser related.

Some network related exceptions are errors that are simple to identify on the server-side by examining the log files. BrowserMob was used as an embedded proxy to intercept all communications between the browser and the server. This allows for the logging of network related exceptions, as well as response interception and subsequent modification for JavaScript exception reporting.

JavaScript exceptions in modern browsers typically manifest themselves via the browser console object. The console is typically not visible to the end user and needs to be enabled via menu options or browser plug-ins such as Firebug for Firefox. A sample of thrown exception is shown in Figure 3.3.

Figure 3.3: An example JavaScript error in the Firebug console.

Unfortunately, once these exceptions are thrown, there is no way to retrieve them from the browser. A simple way to capture these is to introduce a small snippet of JavaScript code in every HTML entry point that has a collection for inserting exceptions. This requires that testing code be added to the deployable deliverable, considered by many to be bad practice. As an alternative, this experiment introduces a novel approach to keep testing scaffolding out of the deployable code base. BrowserMob is used in to intercept only full HTML page responses from the server to the browser, and alter it injecting a custom browser Console object at the beginning of the HTML <script> tag. This custom Console object stores any exception for later retrieval and reporting. Figure 3.4 shows the custom JavaScript injected into each response.

```

100 <script type="text/javascript">
101     window.jsErrors = [];
102     window.onerror = function(errorMessage) {
103         window.jsErrors [window.jsErrors .length] = errorMessage;
104     }
105 </script>

```

Figure 3.4: Browser exception catching

3.8 Combinatorial Testing

The preparation done with the creation of both the exhaustive and the reduced navigation graph has prepared for the generation of combinatorial test cases. The sample application has some variables, namely state, county, and city, that are hierarchical in nature. Czerwonka [10], while describing the capabilities of the Microsoft PICT combinatorial test generation tool, explains that these hierarchical variables are treated as a “sub-model” and pairwise combined first to represent a single variable that is then used in the generation of the combinatorial test cases. For example, in the sample application on the home page, the State of Utah, Cache County, and Logan City are combined together to form a single variable that is then used in building the combinatorial tests. Additionally, a great majority of exceptions that occur in rich web applications result from the different browsers and environments. Tatsumi [11] made the distinction of input and environment variables. The input variables discovered via the process introduced previously can then be combined with provided environment variables to also identify “compatibility faults.” I discuss in the concluding chapter future work that can add distributed computing capabilities enabling browsers in different client environments to perform the same tests.

CHAPTER 4

RESULTS

Each section of this chapter focuses on the results of one of the four research questions discussed in the experiments chapter. First, an abstract URL test suite is compared to the full exhaustive test suite in both size and effectiveness. Second, a more effective combinatorial coverage based test suite is compared to the same exhaustive test suite.

4.1 Size Impact of Abstract URL Test Suite

Using the algorithm described in Figure 2.1, one is able to identify variables in semantic URLs. These variables are then used to build abstract URLs that represent a single application state for each template or combinations of templates in the sample application. The abstract URL test suite does not represent every possible data combination represented in the sample application. The exhaustive count represents every application state as a combination of data plus templates. The hypothesis of this experiment was that the abstract URL test suite would find less faults than the exhaustive test suite but would miss most, if not all, data related faults.

For this relatively small application, one can quickly understand the value of identifying abstract URLs. Indeed, assuming an average page load latency of around 500 milliseconds, running the exhaustive test suite would require more than 27 hours to perform a full regression test.

One option for future work would be to increase the number of abstract URL tests from a single random data point for each application state to a statistically relevant sample of random data points for each application state and compare the effectiveness to the exhaustive test suite. Table 4.1 compares the sizes of the two strategies.

Table 4.1: Size of Test Suites by Strategy.

Strategy	Number of Tests
Exhaustive	199,484
Abstract URLs	27

4.2 Effectiveness of Abstract URL Test Suite

The exhaustive test suite yields significantly better results in fault detection compared to the abstract URL test suite but also took the longest amount of time. In real-world scenarios, it may not be practical to let a full regression test of nearly 200,000 application states run for more than a day before deploying.

Table 4.2: Effectiveness of Single Coverage With Abstract URL Reduction.

Strategy	Number Faults Identified	Percent Faults Identified
Exhaustive	69	89.6%
Abstract URL	15	19.5%

4.3 Size Impact of Combinatorial Coverage

The test generation code, when discovering a `<form>` element on any snippet of HTML code, analyzes it looking for related or hierarchical variables to apply combinatorial algorithms to test generation. Not all of the application states in the sample application have multiple variable combinations as an option. Some pages, such as the application “about” page, have no variables. As the combinatorial test generation algorithms do not apply to them, I have combined both generated combinatorial tests and the remaining “uncombined” states (i.e., about page) to provide a more comprehensive test suite to better compare with the exhaustive, single coverage testing strategy. The single, exhaustive coverage and combinatorial testing coverage had the following number of tests in their respective test suites.

Table 4.3: Size of Test Suites by Strategy.

Strategy	Number of Tests
Exhaustive, Single Coverage	199,484
2-way Combinatorial Coverage	25,560 (13,387 combinatorial + 12,173 remaining)

```

100 <fieldset data-name="unit" data-role="controlgroup">
101   <select data-hierarchical="unit" name="state" id="state">
102     <option>Select State</option>
103   </select>
104   <select data-hierarchical="unit" name="county" id="county">
105     <option value="">Select County</option>
106   </select>
107   <select data-hierarchical="unit" name="city" id="city">
108     <option value="">Select City</option>
109   </select>
110 </fieldset>

```

Figure 4.1: Hierarchical variable HTML form example

Inside an encountered `<form>` element, the combinatorial test generation will search for any input element or select element as well as a “data-hierarchical” attribute that indicates these options should be pre-combined before applied as a option in a variable in the combinatorial test generation. An example `<form>` element can be found in Figure 4.1. The “data-hierarchical” attribute tells the test generation code to combine them as a sub-model before applying it to a combinatorial algorithm. A sample of the pre-combined sub-model data would look like is found in Table 4.4.

With the unit variable type pre-combined in the sample application, combinatorial test generation can then be applied according to Table 4.4.

Table 4.4: Sample Combinatorial Tests.

Unit	Year	Data Point
/state/alabama	1990	Population
/state/alabama/county/autauga	2000	Density
/state/alabama/county/lee/city/auburn-city	2005	Ranking
/state/alabama/county/tuscaloosa/city/tuscaloosa-city		
...		
Remaining combined units		

4.4 Effectiveness of Combinatorial Coverage

The single, exhaustive coverage testing yielded a slightly better fault detection but required a full order of magnitude more tests to achieve that result, 199,484 compared to 25,560. Assuming an average latency of 500 milliseconds would require only 3.5 hours to execute all test cases in the combinatorial coverage test suite, or approximately 23.5 hours less time required to test the exhaustive test suite on the sample application.

Table 4.5: Effectiveness of Combinatorial Coverage.

Strategy	Number Faults Identified	Percent Faults Identified
Exhaustive, Single Coverage	69	89.6%
2-way, Combinatorial Coverage	68	88.3%

CHAPTER 5

CONCLUSIONS

Combinatorial testing is a powerful tool in identifying software faults. Indeed, Kuhn, et al. showed in a study of a NASA Distributed Database that 93% of all faults were identified by 2-way combinations, and 98% by 3-way combinations. Across industries, “the detection rate curves for the other applications studied are similar, reaching 100% detection with 4 to 6-way interactions” [12].

Rich web applications continue to grow with the introduction and rapid development of new HTML5 features and APIs, powerful JavaScript based frameworks, and increasingly more powerful client machines. This thesis demonstrates that combinatorial testing can play an important role in the testing of rich web applications and that more future work is needed.

While there is not a specific standard for semantic URLs compared to the traditional URL format defined in RFC 1378, this paper has shown a novel way to identify variables by employing graph theory and branching complexity analysis. This approach worked with the sample application but needs additional work to be universally applicable to all semantic URL formats. For instance, an application that has a small number of variables that equal the branching complexity of the URL structure could result in false positives with variable identification.

This research also showed that using abstract URLs to generate test cases was an effective and inexpensive way to discover all types of faults except data faults in the sample application and particularly well suited for template heavy client-side applications. While it didn’t find as many errors as the exhaustive or combinatorial approaches, with only 27 tests (0.0001% of the exhaustive tests run) it found 19.5%. Using abstract URLs may be a good strategy when time is extremely limited. A point of future work that may possibly yield

better results with abstract URLs would be to increase the number of random variables used from a single variable to a statistically significant percentage of total data points available.

Also demonstrated in this paper was a convenient way to capture JavaScript exceptions by intercepting HTTP requests and responses via an embedded proxy server, and then injecting a JavaScript array in the <HEAD> section of each HTML page to capture any thrown exceptions.

Future work may extend testing to distributed machines and different client environments. For example, the environment variables shown in Table 5.1 could be combined with input variables to help catch compatibility faults, an ongoing concern with the various browser creators adoption rate of new HTML5 features.

Table 5.1: Environment Variables.

OS	Browser
Linux	Chrome
OSX	Firefox
Windows	IE
	Safari

While I attempted to make the sample application as “real world” as possible, additional work is needed to make the testing tool better and more practical for rich web applications in real world scenarios. An initial effort was made to implement ideas in Microsoft’s PICT tool [10], such as making a better distinction between preparation and test generation, pre-combining related, hierarchical fields, and providing test generation guidelines by employing the new HTML5 data-* attributes. As an example of future enhancements to the testing software, one could create a new “data-exclude-test” to keep a particular variable from being combined for test case generation. Also, guidelines on associated form variables with attributes in the `<fieldset>` tags would allow for better control of test case generation with variables. Integrating in with build tools would also be beneficial to keep those guideline attributes from being deployed to production. More future work that would be beneficial would be better testing of all JavaScript event types and application state

transitions that are not associated with an update to the URL fragment (i.e., an event triggered by the clicking of a non anchor tag that updates a value in an existing portion of the Document Object Model (DOM) structure).

Despite much future work left to be researched, this thesis demonstrates that test case generation using combinatorial coverage strategies in rich web applications, as in other application types, provides much benefit in identifying faults and should be explored further.

REFERENCES

- [1] T. Berners-Lee, L. Masinter, M. McCahill *et al.*, “Uniform resource locators (URL),” <http://www.ietf.org/rfc/rfc1738.txt>, 1994, [Online; accessed June 8, 2012].
- [2] T. Burners-Lee, “Cool URIs don’t change,” *W3C*, 2008.
- [3] A. Gleyzer, M. Denisyuk, A. Rimmer, and Y. Salingar, “A fast recursive GIS algorithm for computing strahler stream order in braided and nonbraided networks,” *JAWRA Journal of the American Water Resources Association*, vol. 40, no. 4, pp. 937–946, 2004.
- [4] S. Sampath, S. Sprenkle, E. Gibson, L. Pollock, and A. Greenwald, “Applying concept analysis to user-session-based testing of web applications,” *Software Engineering, IEEE Transactions on*, vol. 33, no. 10, pp. 643–658, 2007.
- [5] Y. Guo and S. Sampath, “Web application fault classification—an exploratory study,” in *Proceedings of the Second ACM-IEEE international symposium on Empirical software engineering and measurement*. ACM, 2008, pp. 303–305.
- [6] F. Ocariza Jr, K. Pattabiraman, and B. Zorn, “JavaScript errors in the wild: An empirical study,” in *Software Reliability Engineering (ISSRE), 2011 IEEE 22nd International Symposium on*. IEEE, 2011, pp. 100–109.
- [7] Mozilla, “Mozilla Developer Network, Error Object Reference,” https://developer.mozilla.org/en/JavaScript/Reference/Global_Objects/Error, 2011, [Online; accessed June 8, 2012].

- [8] I. Herman, G. Melançon, and M. Marshall, “Graph visualization and navigation in information visualization: A survey,” *Visualization and Computer Graphics, IEEE Transactions on*, vol. 6, no. 1, pp. 24–43, 2000.
- [9] W. Wang, Y. Lei, S. Sampath, R. Kacker, R. Kuhn, and J. Lawrence, “A combinatorial approach to building navigation graphs for dynamic web applications,” in *Software Maintenance, 2009. ICSM 2009. IEEE International Conference on*. IEEE, 2009, pp. 211–220.
- [10] J. Czerwonka, “Pairwise testing in the real world: Practical extensions to test-case scenarios,” <http://msdn.microsoft.com/en-us/library/cc150619.aspx>, 2008, [Online; accessed June 8, 2012].
- [11] M. Grochtmann, J. Wegener, and K. Grimm, “Test case design using classification trees and the classification-tree editor CTE,” in *Proceedings of Quality Week*, vol. 95, 1995, p. 30.
- [12] D. Kuhn, R. Kacker, and Y. Lei, “Practical combinatorial testing,” *NIST Special Publication*, vol. 800, p. 142, 2010.