

ESGALDNS:
NAMECOIN-BASED ANONYMOUS DOMAIN NAME SERVICE
FOR TOR HIDDEN SERVICES

by

Jesse Victors

A thesis submitted in partial fulfillment
of the requirements for the degree

of

MASTER OF SCIENCE

in

Computer Science

Approved:

Dr. Ming Li
Major Professor

Dr. Nicholas Flann
Committee Member

Dr. Daniel Watson
Committee Member

Dr. Mark R. McLellan
Vice President for Research and
Dean of the School of Graduate Studies

UTAH STATE UNIVERSITY
Logan, Utah

2015

Copyright © Jesse Victors 2015

All Rights Reserved

ABSTRACT

EsgalDNS:
Namecoin-based Anonymous Domain Name Service
for Tor Hidden Services

by

Jesse Victors, Master of Science
Utah State University, 2015

Major Professor: Dr. Ming Li
Department: Computer Science

The Tor network is a second-generation onion routing system that aims to provide anonymity, privacy, and Internet censorship resistance to its users. In recent years it has grown significantly in response to revelations of national and global electronic surveillance, and remains one of the most popular and secure anonymity network in use today. Tor is also known for its support of anonymous websites within its network. Decentralized and secure, the domain names for these services are tied to public key infrastructure (PKI) but are challenged by their long and technical addresses. In response to this difficulty, in this thesis I introduce a novel and decentralized Tor-powered DNS system that provides unique and human-meaningful domain names to Tor hidden services.

(28 pages)

PUBLIC ABSTRACT

Jesse M. Victors

The Tor network is a second-generation onion routing system that aims to provide anonymity, privacy, and Internet censorship resistance to its users. In recent years it has grown significantly in response to revelations of national and global electronic surveillance, and remains one of the most popular and secure anonymity network in use today. Tor is also known for its support of anonymous websites within its network. Decentralized and secure, the domain names for these services are tied to public key infrastructure (PKI) but are challenged by their long and technical addresses. In response to this difficulty, in this thesis I introduce a novel and decentralized Tor-powered DNS system that provides unique and human-meaningful domain names to Tor hidden services.

This work is dedicated to the developers and community behind the Tor Project and the Tails OS.

CONTENTS

	Page
ABSTRACT	iii
PUBLIC ABSTRACT	iv
LIST OF TABLES	vii
LIST OF FIGURES	viii
CHAPTER	
1 INTRODUCTION	1
2 BACKGROUND	2
2.1 Tor	2
2.2 Bitcoin	7
2.3 Namecoin	9
2.4 DNS	10
2.5 Zooko's Triangle	10
2.6 Summary	11
3 REQUIREMENTS	12
4 CHALLENGES	13
5 EXISTING WORKS	14
6 SOLUTION	15
7 ANALYSIS	16
8 RESULTS	17
9 FUTURE WORK	18
10 CONCLUSION	19
REFERENCES	20

LIST OF TABLES

Table

Page

LIST OF FIGURES

Figure		Page
2.1	Anatomy of the construction of a Tor circuit.	4
2.2	A circuit through the Tor network.	4
2.3	A Tor circuit is changed periodically, creating a new user identity.	5
2.4	Alice uses the encrypted cookie to tell Bob to switch to <i>RP</i>	6
2.5	Bidirectional communication between Alice and the hidden service.	6
2.6	A sample blockchain.	8
2.7	Three traditional Bitcoin transactions.	9
2.8	Zooko's Triangle.	11

CHAPTER 1

INTRODUCTION

The Tor network is a second-generation onion routing system that aims to provide anonymity, privacy, and Internet censorship protection to its users. The Tor client software multiplexes all end-user TCP traffic through a series of relays on the Tor network, typically a carefully-constructed three-hop path known as a *circuit*. Each relay in the circuit has its own encryption layer, so traffic is encrypted multiple times and then is decrypted in an onion-like fashion as it travels through the Tor circuit. As each relay sees no more than one hop in the circuit, in theory neither an eavesdropper nor a compromised relay can link the connection's source, destination, and content. Tor remains one of the most popular and secure tools to use against network surveillance, traffic analysis, and information censorship.

While the majority of Tor's usage is for traditional access to the Internet, Tor's routing scheme also supports anonymous websites, hidden inside Tor. Unlike the Clearnet, Tor does not contain a traditional DNS system for its websites; instead, hidden services are identified by their public key and can be accessed through Tor circuits. A client and the hidden service can thus communicate anonymously.

CHAPTER 2

BACKGROUND

This chapter is divided into five main sections. First, I provide background on the Tor anonymity network and describe how clients can contact hidden services inside it. Secondly, I explain Bitcoin and its architecture as a prerequisite for Namecoin. Third, I detail Namecoin and compare and contrast with Bitcoin. Fourth, I provide an overview of traditional DNS and how it is used on the clearnet, and finally I introduce Zooko's Triangle: a conjecture that makes claims on the limits of persistent name systems.

2.1 Tor

The Tor network is a second-generation onion routing system that aims to provide anonymity, privacy, and Internet censorship protection to its users. Tor routes encrypted TCP/IP user traffic through a worldwide volunteer-run network of over six thousand relays. Tor's encryption, authentication, and routing protocols are designed to make it very difficult for any adversary to identify an end user or correlate them to their traffic. Tor continues to be one of the most popular and secure tools to use against network surveillance, traffic analysis, and information censorship.

2.1.1 Design

Tor provides an anonymity and privacy layer by relaying all end-user TCP traffic through a series of relays on the Tor network. Typically this route consists of a carefully-constructed three-hop path known as a *circuit*, which changes over time. These nodes in the circuit are commonly referred to as *guard node*, *middle relay*, and the *exit node*, respectively. Only the first node is exposed to the origin of TCP traffic into Tor, and only the exit node can see the destination of traffic out of Tor. The middle router, which passes encrypted

traffic between the two, is unable to determine either. As such, each node is only aware of the machines it talks to, and only the client knows the identity of all three nodes used in its circuit. Tor's architecture thus minimizes the security risk from compromised nodes. [1]

Tor also contains nine authority nodes that maintain a list of IPs, ports, public keys, status, capability flags (entry guard, exit, etc) and other information on all Tor nodes in the network. This list is then signed by these nodes and periodically copied across all Tor nodes so that all Tor nodes can contact and authenticate to all other Tor nodes in the network.

2.1.2 Routing

In traditional Internet connections, the client communicates directly with the server. In this model, an eavesdropper can often reveal both the identity of the end user and their activities. Direct encrypted connections do not hide IP headers, which expose source and destination addresses and the size of the payload. In the face of adversaries with sophisticated traffic analysis tools, such information can be very revealing for someone who wishes to hide their online activities.

Tor combats this by routing end user traffic through a randomized circuit through the network of relays. The Tor client software first queries a trusted directory server or a relay mirroring the directory. This directory contains a list of IPs, ports, public keys, and other information about all nodes in the Tor network. [2] Next, the Tor client chooses three unique and geographically diverse nodes to use. It then builds and extends the circuit one node at a time, negotiating respective HTTPS connections with each node in turn. No single relay knows the complete path, and each relay can only decrypt its layer of decryption. In this way, data is encrypted multiple times and then is decrypted in an onion-like fashion as it passes through the circuit.

The client first establishes a TLS connection with the first relay, R_1 , using the relay's public key. The client then performs a Diffie-Hellman-Merkle key exchange to negotiate K_1 which is then used to generate two symmetric session keys: a forward key $K_{1,F}$ and a backwards key $K_{1,B}$. $K_{1,F}$ is used to encrypt all communication from the client to R_1 and $K_{1,B}$ is used for all replies from R_1 to the client. These keys are used in conjunction with the

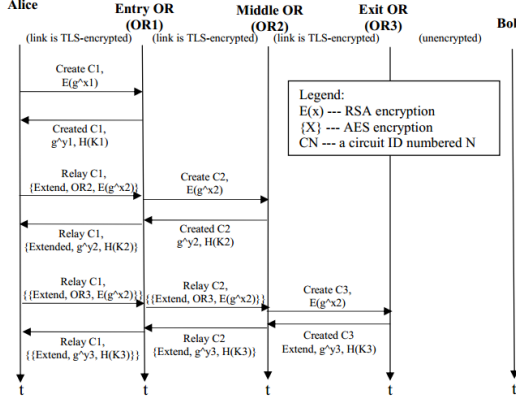


Figure 2.1: Anatomy of the construction of a Tor circuit.

How Tor Works




Figure 2.2: A circuit through the Tor network.

symmetric cipher suite negotiated during the TLS handshake, thus forming an encrypted tunnel with perfect forward secrecy. Once this one-hop circuit has been created, the client then sends R_1 the RELAY_EXTEND command, the address of R_2 , and the client's half of the Diffie-Hellman-Merkle protocol using $K_{1,F}$. R_1 performs a TLS handshake with R_2 and uses R_2 's public key to send this half of the handshake to R_2 , who replies with his second half of the handshake and a hash of K_2 . R_1 then forwards this to the client under $R_{1,B}$ with the RELAY_EXTENDED command to notify the client. The client generates $K_{1,F}$ and $K_{1,B}$ from K_2 , and repeats the process for R_3 , [3] as shown in Figure 3. The TLS/IP connections remain open, so the returned information travels back up the circuit to the end user.

Following the complete establishment of a circuit, the Tor client software then offers a Secure Sockets (SOCKS) interface on localhost which multiplexes TCP traffic through Tor. At the application layer, this data is packed and padded into equally-sized Tor *cells*, transmission units of 512 bytes. As each relay sees no more than one hop in the circuit, in theory neither an eavesdropper nor a compromised relay can link the connection's source, destination, and content. Tor further obfuscates user traffic by changing the circuit path every ten minutes, [1] as shown in Figure 4. A new circuit can also be requested manually by the user.

If the recipient or web server supports encryption, often a HTTPS or traditional

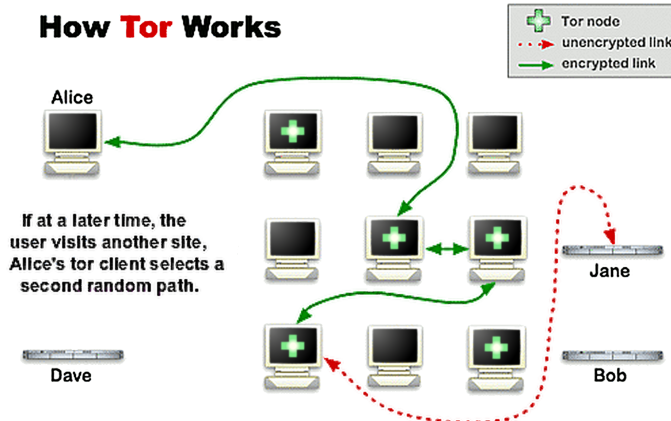


Figure 2.3: A Tor circuit is changed periodically, creating a new user identity.

TLS connection will be established with the other client or web server. If this happens, end-to-end encryption is complete and the exit node cannot see the user's traffic, and an outsider near the user would be faced with up to four layers of TLS encryption: $K_{1,F}(K_{2,F}(K_{3,F}(K_{server}(\text{client request}))))$ and likewise $K_{1,B}(K_{2,B}(K_{3,B}(K_{server}(\text{server reply}))))$ for the returning traffic. This makes traffic analysis and cryptographic attacks very difficult.

Tor users typically use the Tor Browser Bundle, (TBB) a custom build of Mozilla Firefox with a focus on security and privacy. The TBB anonymizes and provides privacy to the user in many ways. These include blocking all web scripts not explicitly whitelisted, forcing all traffic including DNS requests through the Tor SOCKS port, mimicking Firefox in Windows both with a user agent (regardless of the native platform) and SSL cipher suites, and reducing Javascript timer precision to avoid identification through clock skew. Furthermore, the TBB includes the Electronic Frontier Foundation's HTTPS Everywhere extension, which uses regular expressions to rewrite HTTP web requests into HTTPS whenever possible. Thus, if the web server is capable of handling SSL or TLS connections, HTTP communications will be encrypted to them. If this is the case, the TBB performs a TLS handshake with the web server, but the exchange happens through the Tor circuit. This provides the final layer of encryption to the outside.

2.1.3 Hidden Services

While the majority of Tor's usage is for traditional access to the Internet, Tor's routing scheme also supports anonymous services, such as websites, marketplaces, or chatrooms. These are a part of the Deep Web and cannot be normally contacted outside of Tor. Unlike the Clearnet, Tor does not contain a traditional DNS system for its websites; instead, every hidden service has a public and private RSA key, and domains are a truncated SHA-1 hash of its public key. This means that domain names are distributed and correlate directly to the identity of a hidden service, allowing anyone to verify the authenticity of the service server, akin to SSL certificates on the Clearnet. Tor hidden services allow a client, Alice, and a hidden service, Bob, to communicate anonymously. [?]

In advance, Bob builds Tor circuits to several random relays and enables them to act as *introduction points* by giving them his public key, B_K . He then uploads this information to a distributed hashtable inside the Tor network, signing the result. Alice queries this hashtable, finds B_K and his introduction points, and builds a Tor circuit to one of them, IP_1 . Simultaneously, she also builds a circuit to another relay, RP , which she enables as a rendezvous point by telling it a one-time secret, S .




Figure 2.4: Alice uses the encrypted cookie to tell Bob to switch to RP .

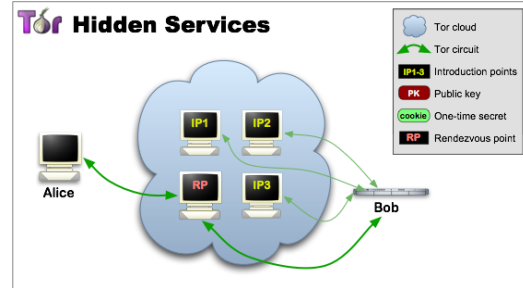


Figure 2.5: Bidirectional communication between Alice and the hidden service.

She then sends to IP_1 an cookie encrypted with B_K , containing RP and S . Bob decrypts this message, builds a circuit to RP , and tells it S_1 , enabling Alice and Bob to communicate. Their communication travels through six Tor nodes: three established by Alice and three by Bob, so both parties remain anonymous. This method for addressing

hidden services has never been revised since the protocol was introduced in 2004. [4]

2.2 Bitcoin

Bitcoin is a decentralized peer-to-peer digital cryptocurrency, created by pseudonymous developer Satoshi Nakamoto in 2008. Ownership of Bitcoins consists of holding a private ECDSA key, and a transfer is a transmission of Bitcoins from one key to another. All transactions are recorded on a public ledger, called a blockchain, a data structure whose integrity is ensured through computational power but publically verifiable. Bitcoins are generated computationally at a fixed rate by *miners* in a process that also secures the blockchain. Although Bitcoin received limited attention in the first two years of its life, it has since grown significantly since then, with approximately 70,000 daily transactions as of the time of this writing. Bitcoin's growth has led to the creation of many alternative cryptocurrencies, and its popularity has influenced financial discussions and legal controversy worldwide.

2.2.1 Architecture

A blockchain is data structure fundamental to Bitcoin, and crucial for its functionality. As a distributed decentralized system, this public ledger is Nakamoto's answer to the problem of ensuring agreement of critical data across all involved parties. The blockchain is a novel structure, and its structure guarantees integrity, chronological ordering of transactions, and the prevention of double-spending of Bitcoins. The blockchain consists of blocks of data that are held together by proof-of-work, a cryptographic puzzle whose solution is provably hard to find but trivial to verify. Bitcoin's proof-of-work is based on Adam Back's Hashcash scheme: that is, find a nonce such that the hash of this nonce and some data produces a result that begins with a certain number of zero bits. In Bitcoin's case this is stated as finding a nonce that when passed through two rounds of SHA256 (SHA256²) produces a value less than or equal to a target T . This requires a party to perform on average $\frac{1}{Pr[H \leq T]} = \frac{2^{256}}{T}$ amount of computations, but it is easy to verify that $SHA256^2(msg||n) \leq T$. Nodes in the Bitcoin network collectively agree to use the blockchain with the highest accumulation of computational effort, so an adversary seeking to modify the structure would

In the possibility that multiple nodes solve the proof-of-work and generate a new block simultaneously, the block becomes orphaned, the transactions recycled, and the blockchain follows the longest path from the genesis node to the latest block. Each transaction contains the public of the recipient, the ECDSA digital signature of the transaction from the sender, and the hash of the originating transaction. In this way, the digital signatures and proof-of-work in the blockchain can be traced back to the origin and forwards indefinitely.




Figure 2.7: Three traditional Bitcoin transactions.

2.3 Namecoin

Namecoin is a decentralized information registration and transfer system based on Bitcoin. It was the first software fork of Bitcoin and was introduced in April 2011. It uses its own blockchain and can hold name-value pairs in the blockchain attached to coins. While Bitcoin is primarily focused on supporting a currency, Namecoin aims to be a general key-value store, capable of holding cryptographic keys, DNS registrations, or other arbitrary data. It is most commonly used as a secure and censorship-resistance replacement for clearnet DNS. In 2014, Namecoin was recognized by ICANN as the most well-known example of a PKI and DNS system with an emphasis of distributed control and privacy, a growing trend in light of the revelations about the US Government by Edward Snowden.

2.3.1 Names

Although it inheriting Bitcoin's existing infrastructure, Namecoin added several transaction types specifically for registering and processing names, along with two new rules: names in the blockchain expire after 36,000 blocks unless renewed by the owner and no two unexpired names can be identical. These rules are enforced in the blockchain by Namecoin nodes and anyone verifying the Namecoin blockchain. Registering a name consumes 0.01 Namecoin, names can also be transferred to other owners, and they are two types: DNS and personal. The DNS type uses a new Top Level Domain (TLD) not in use by ICANN: .bit, and is used for DNS registrations. The personal name can contain arbitrary data, including

user information such as cryptographic keys. Like Bitcoin, Namecoin’s maximum block size is one megabyte and the difficulty is set such that blocks generate every 10 minutes. Thus names expire every 250 days.

2.4 DNS

The Internet Domain Name Service (DNS) is a hierarchical distributed naming system for computers connected to the Internet. It links two principal Internet namespaces, Internet Protocol (IP) addresses and domain names, and translates one to the other. IP addresses specify the location of a computer or device on a network and domain names identify that resource. Domain names also serve as an abstraction layer so that devices can be moved to a different physical location or to a different IP address without loss of functionality. In contrast to IP addresses, domain names are human-meaningful and easily memorized, so DNS is a crucial component to the usability of the Internet.

Domain names on the Internet consist of a sequence of labels, delimited by dots. The right-most label is the top-level domain (TLD) and can be used to classify the Internet resource by country or by organization type, although generic TLDs are more common. One or more subdomains follow the TLD. Each label can consist of up to 63 characters and the domain names can be up to 253 characters.

2.5 Zooko’s Triangle

Zooko’s Triangle is an influential conjecture proposed by Zooko Wilcox-O’Hearn in late 2001. The conjecture states that in a persistent naming system, only two out of the three following properties can be established: [6]

- Human meaningfulness: the names have a quality of meaningfulness and memorability to the users of the system.
- Securely one-to-one: each name is unique, corresponds to a unique entity or owner, and cannot be forged or mimicked.

- Distributed: the naming system lacks a central authority or database for allocating and distributing names.




Figure 2.8: Zooko’s Triangle.

For example, Tor hidden service .onion addresses and Bitcoin addresses are secure and decentralized but are not human-meaningful. Internet domain names are memorable and provably collision-free, but use central database managed DNS under the jurisdiction of ICANN. Finally, human nicknames are meaningful and distributed, but not securely collision-free. [7]

In January of 2011, Aaron Swartz described a naming system based on Bitcoin that achieved all three properties and thus completed Zooko’s Triangle. Three months later Namecoin was released as a proof-of-concept, becoming the first system to violate the conjecture. As described above, it allowed human-meaningful key-value pairs (names) to be secured in a global blockchain by proof-of-work, and specified that the uniqueness of the names be verified by a distributed network of miners and confirmation nodes.

2.6 Summary

In this thesis, I utilize the Namecoin blockchain to implement a distributed DNS system that provides human-meaningful and provably unique names to Tor hidden services.

CHAPTER 3

REQUIREMENTS

CHAPTER 4

CHALLENGES

CHAPTER 5

EXISTING WORKS

CHAPTER 6

SOLUTION

CHAPTER 7

ANALYSIS

CHAPTER 8

RESULTS

CHAPTER 9

FUTURE WORK

CHAPTER 10

CONCLUSION

Conclusion!

REFERENCES

- [1] D. McCoy, K. Bauer, D. Grunwald, T. Kohno, and D. Sicker, “Shining light in dark places: Understanding the tor network,” *Department of Computer Science and Engineering, University of Washington, Seattle, WA 98195-2969*, 2008.
- [2] L. Xin and W. Neng, “Design improvement for tor against low-cost traffic attack and low-resource routing attack,” *2009 International Conference on Communications and Mobile Computing*, 2009.
- [3] Z. Ling, J. Luo, W. Yu, X. Fuc, W. Jia, and W. Zhao, “Protocol-level attacks against tor,” *Computer Networks*, 2012.
- [4] S. Nicolussi, “Human-readable names for tor hidden services,” 2011.
- [5] K. Okupski, “Bitcoin developer reference,” 2014.
- [6] M. S. Ferdous, A. Jsang, K. Singh, and R. Borgaonkar, “Security usability of petname systems,” October 2009.
- [7] M. Stiegler, “Petname systems,” 2005.